Input0
[forward_compute_time=0.000,backward_compute_time=0.000,activation_size=0.0,parameter_size=0.0]
Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
[forward_compute_time=0.770,backward_compute_time=0.965,activation_size=16056320.0,parameter_size=37632.0]
BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.125,backward_compute_time=0.734,activation_size=16056320.0,parameter_size=512.0]
ReLU(inplace)
[forward_compute_time=1.032,backward_compute_time=0.726,activation_size=16056320.0,parameter_size=0.0]
MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
[forward_compute_time=1.016,backward_compute_time=0.739,activation_size=4014080.0,parameter_size=0.0]
Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.102,backward_compute_time=0.863,activation_size=4014080.0,parameter_size=147456.0]
BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.121,backward_compute_time=0.658,activation_size=4014080.0,parameter_size=512.0]
ReLU(inplace)
[forward_compute_time=0.983,backward_compute_time=0.636,activation_size=4014080.0,parameter_size=0.0]
Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.078,backward_compute_time=0.881,activation_size=4014080.0,parameter_size=147456.0]
BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.103,backward_compute_time=0.661,activation_size=4014080.0,parameter_size=512.0]
Add(inplace)
[forward_compute_time=0.000,backward_compute_time=0.000,activation_size=4014080.0,parameter_size=0.0]
ReLU(inplace)
[forward_compute_time=0.999,backward_compute_time=0.639,activation_size=4014080.0,parameter_size=0.0]
Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.087,backward_compute_time=0.857,activation_size=4014080.0,parameter_size=147456.0]
BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.097,backward_compute_time=0.686,activation_size=4014080.0,parameter_size=512.0]
ReLU(inplace)
[forward_compute_time=1.013,backward_compute_time=0.617,activation_size=4014080.0,parameter_size=0.0]
Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.113,backward_compute_time=0.945,activation_size=4014080.0,parameter_size=147456.0]
BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.082,backward_compute_time=0.601,activation_size=4014080.0,parameter_size=512.0]
Add(inplace)
[forward_compute_time=0.000,backward_compute_time=0.000,activation_size=4014080.0,parameter_size=0.0]
ReLU(inplace)
[forward_compute_time=0.984,backward_compute_time=0.567,activation_size=4014080.0,parameter_size=0.0]
Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
[forward_compute_time=1.094,backward_compute_time=0.844,activation_size=2007040.0,parameter_size=294912.0]
BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.069,backward_compute_time=0.551,activation_size=2007040.0,parameter_size=1024.0]
Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
[forward_compute_time=1.032,backward_compute_time=0.711,activation_size=2007040.0,parameter_size=32768.0]
ReLU(inplace)
[forward_compute_time=1.009,backward_compute_time=0.538,activation_size=2007040.0,parameter_size=0.0]
BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.061,backward_compute_time=0.534,activation_size=2007040.0,parameter_size=1024.0]
Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.106,backward_compute_time=0.707,activation_size=2007040.0,parameter_size=589824.0]
BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.077,backward_compute_time=0.559,activation_size=2007040.0,parameter_size=1024.0]
Add(inplace)
[forward_compute_time=0.000,backward_compute_time=0.000,activation_size=2007040.0,parameter_size=0.0]
ReLU(inplace)
[forward_compute_time=1.002,backward_compute_time=0.504,activation_size=2007040.0,parameter_size=0.0]
Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.087,backward_compute_time=0.676,activation_size=2007040.0,parameter_size=589824.0]
BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.092,backward_compute_time=0.620,activation_size=2007040.0,parameter_size=1024.0]
ReLU(inplace)
[forward_compute_time=0.995,backward_compute_time=0.490,activation_size=2007040.0,parameter_size=0.0]
Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.095,backward_compute_time=0.780,activation_size=2007040.0,parameter_size=589824.0]
BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.060,backward_compute_time=0.471,activation_size=2007040.0,parameter_size=1024.0]
Add(inplace)
[forward_compute_time=0.000,backward_compute_time=0.000,activation_size=2007040.0,parameter_size=0.0]
ReLU(inplace)
[forward_compute_time=0.992,backward_compute_time=0.428,activation_size=2007040.0,parameter_size=0.0]
Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
[forward_compute_time=1.122,backward_compute_time=0.787,activation_size=1003520.0,parameter_size=1179648.0]
Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
[forward_compute_time=1.010,backward_compute_time=0.804,activation_size=1003520.0,parameter_size=131072.0]
BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.060,backward_compute_time=0.475,activation_size=1003520.0,parameter_size=2048.0]
BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.015,backward_compute_time=0.511,activation_size=1003520.0,parameter_size=2048.0]
ReLU(inplace)
[forward_compute_time=0.988,backward_compute_time=0.491,activation_size=1003520.0,parameter_size=0.0]
Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.132,backward_compute_time=0.708,activation_size=1003520.0,parameter_size=2359296.0]
BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.037,backward_compute_time=0.494,activation_size=1003520.0,parameter_size=2048.0]
Add(inplace)
[forward_compute_time=0.000,backward_compute_time=0.000,activation_size=1003520.0,parameter_size=0.0]
ReLU(inplace)
[forward_compute_time=0.966,backward_compute_time=0.503,activation_size=1003520.0,parameter_size=0.0]
Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.134,backward_compute_time=0.680,activation_size=1003520.0,parameter_size=2359296.0]
BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.043,backward_compute_time=0.541,activation_size=1003520.0,parameter_size=2048.0]
ReLU(inplace)
[forward_compute_time=0.983,backward_compute_time=0.511,activation_size=1003520.0,parameter_size=0.0]
Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.087,backward_compute_time=1.159,activation_size=1003520.0,parameter_size=2359296.0]
BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.026,backward_compute_time=0.554,activation_size=1003520.0,parameter_size=2048.0]
Add(inplace)
[forward_compute_time=0.000,backward_compute_time=0.000,activation_size=1003520.0,parameter_size=0.0]
ReLU(inplace)
[forward_compute_time=0.943,backward_compute_time=0.575,activation_size=1003520.0,parameter_size=0.0]
Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
[forward_compute_time=1.111,backward_compute_time=1.085,activation_size=501760.0,parameter_size=4718592.0]
Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
[forward_compute_time=0.919,backward_compute_time=1.141,activation_size=501760.0,parameter_size=524288.0]
BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=1.001,backward_compute_time=0.544,activation_size=501760.0,parameter_size=4096.0]
BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=0.919,backward_compute_time=0.570,activation_size=501760.0,parameter_size=4096.0]
ReLU(inplace)
[forward_compute_time=0.904,backward_compute_time=0.611,activation_size=501760.0,parameter_size=0.0]
Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.120,backward_compute_time=0.993,activation_size=501760.0,parameter_size=9437184.0]
BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=0.990,backward_compute_time=0.542,activation_size=501760.0,parameter_size=4096.0]
Add(inplace)
[forward_compute_time=0.000,backward_compute_time=0.000,activation_size=501760.0,parameter_size=0.0]
ReLU(inplace)
[forward_compute_time=0.853,backward_compute_time=0.672,activation_size=501760.0,parameter_size=0.0]
Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.095,backward_compute_time=1.069,activation_size=501760.0,parameter_size=9437184.0]
BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=0.936,backward_compute_time=0.607,activation_size=501760.0,parameter_size=4096.0]
ReLU(inplace)
[forward_compute_time=0.816,backward_compute_time=0.682,activation_size=501760.0,parameter_size=0.0]
Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
[forward_compute_time=1.054,backward_compute_time=1.599,activation_size=501760.0,parameter_size=9437184.0]
BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
[forward_compute_time=0.847,backward_compute_time=0.698,activation_size=501760.0,parameter_size=4096.0]
Add(inplace)
[forward_compute_time=0.000,backward_compute_time=0.000,activation_size=501760.0,parameter_size=0.0]
ReLU(inplace)
[forward_compute_time=0.696,backward_compute_time=0.670,activation_size=501760.0,parameter_size=0.0]
AvgPool2d(kernel_size=7, stride=1, padding=0)
[forward_compute_time=0.709,backward_compute_time=0.674,activation_size=10240.0,parameter_size=0.0]
Size(0)
[forward_compute_time=0.000,backward_compute_time=0.000,activation_size=4,parameter_size=0.0]
View(-1)
[forward_compute_time=0.000,backward_compute_time=0.000,activation_size=10240.0,parameter_size=0.0]
Linear(in_features=512, out_features=1000, bias=True)
[forward_compute_time=0.703,backward_compute_time=0.671,activation_size=20000.0,parameter_size=2052000.0]